

# Introduction to Matching Methods for Causal Inference

Zeren Li

## Introduction

---

- Useful for observational studies
- Sometimes better than simply using control variables.
- It's popular! (Since the 1970s, 93,700 articles uses matching method)

## Regression v.s. Matching ¶

---

Both methods assumes **Selection on Observables**

- assume Potential outcomes ( $y_t, y_c$ ) are independent of treatment after controlling for observed covariates  $x$  (Conditional Independence Assumption)
- use a set of covariates to adjust the estimation

Both methods are **weight scheme**. Regression aims to minimize squared errors, so observations on the margins have more weight. Matching puts the emphasis on observations that have similar covariates.

**Two methods can work together! (e.g. Regression with Propensity Score Matching)**

## Why matching

---

**Reduce Model Dependence:**

- Does not rely on modeling assumptions of linear dependencies across the entire range of the control variables to do controlled comparisons.
- Enhances apples to apples comparisons: Making sure the treated and untreated observations are balanced.

## Matching Reduces Model Dependence

---


Figure from Ho, Imai, King, Stuart (2007)

 alt text

## Matching Reduces Model Dependence

---

Figure from Ho, Imai, King, Stuart (2007)

 alt text

## Matching Reduces Model Dependence

---

Figure from Ho, Imai, King, Stuart (2007)

 alt text

## Matching Reduces Model Dependence

---

Figure from Ho, Imai, King, Stuart (2007)

 alt text

## Matching Reduces Model Dependence

---

Figure from Ho, Imai, King, Stuart (2007)

 alt text

## Matching & Common Support

---

- Common support ensures the treated have the control with “close”  $P(x)$
- Lack of common support appears in tails of distributions
- Larger sample of eligible nonparticipants helps matching
- Poor common support can induce bias in matching estimator
  - E.g., if no matches may drop nonrandom subset of participants

## Matching & Common Support

 Khandler et al(2010)

## When Matching

---

- usually used in observational studies
- treatment and control groups are imbalanced
- a lot of covariates may be correlated with Y and D
- pre-process the data before running regression
- as a robustness check

## Steps of Matching

---

### Defining closeness I: Variables to Include

Defining “closeness”: the distance measure used to determine whether an individual is a good match for another

- Be liberal in terms of including variables that may be associated with treatment assignment & outcomes
- In small sample, select variables without using the observed outcomes, but based on the previous research
- DO NOT include post-treatment variables

### Defining closeness II: Metric of Closeness

- Exact
- Mahalanobis
- Propensity Score

## Matching Methods

Implement a matching method, given that measure of closeness

- Nearest neighbor matching
- Subclassification, Full Matching, and Weighting

## Assessing Common Support

- show distributions of propensity scores of treat&control groups
- substantial overlap of the propensity score distributions in the two groups

## Matching Diagnostics

Assess the quality of the resulting matched samples

- diagnose the quality of the matching through an assessment of **covariate balance**
- **covariate balance** is defined as similarity of empirical distributions of the full set of covariates in the matched treated and control groups
- **Numerical diagnostics** (e.g. standardized difference in means): difference in means of each covariate, divided by the standard deviation in the full treated group:  $\frac{X_t - X_c}{\sigma_t}$
- Graphical diagnostics

## Analysis of the outcome

Analyze of the outcome and estimation of the treatment effect

- **Matching** and **Regression** have been shown to work best in combination
- After k:1 matching
  - do same analysis as what did with original data, but using the matched data instead, e.g. regression
  - weight needs to be incorporated
- After subclassification or full matching

## Propensity Score Matching Example

We practise the matching method by replicating the prominent paper on job training. The 'lalonge.dta' consists of the real earnings in the year 1978 (the response), a treatment indicator, and a number of demographic variables (controls). Here are the variable definitions.

- age: age in years.
- educ: years of schooling.
- black: indicator variable for blacks.
- hisp: indicator variable for Hispanics.
- married: indicator variable for marital status.
- nodegr: indicator variable for high school diploma.
- re74: real earnings in 1974.
- re75: real earnings in 1975.
- re78: real earnings in 1978.
- treat: an indicator variable for treatment status.

The original paper is here: Robert Lalonde, "Evaluating the Econometric Evaluations of Training Programs", American Economic Review, Vol. 76, pp. 604-620

```
In [41]: # load packages
# now pymatch only works with older versions of pandas, please make sure
# you have a version older than 0.23.4
# If not, try this to downgrade pandas: pip install pandas==0.23.4
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from pymatch.Matcher import Matcher
%matplotlib inline
```

```
In [42]: lalonge = pd.read_stata("lalonge.dta")
lalonge.head()
```

Out[42]:

	treated	age	education	black	married	nodegree	re74	re75	re78
0	1.0	33	12	0	1	0	0.000000	0.000000	12418.070312
1	1.0	20	12	0	1	0	8644.156250	8644.156250	11656.505859
2	0.0	39	12	1	1	0	19785.320312	6608.137207	499.257202
3	1.0	49	8	0	1	1	9714.596680	7285.947754	16717.121094
4	0.0	26	8	0	1	1	37211.757812	36941.265625	30247.500000

```
In [43]: lalonde.groupby("treated").mean()
```

```
Out[43]:
```

	age	education	black	married	nodegree	re74	re75	r
treated								
0.0	24.447059	10.188235	0.800000	0.157647	0.814118	3672.485151	3026.682738	5090.048
1.0	24.626263	10.380471	0.801347	0.168350	0.730640	3570.998967	3066.098189	5976.352

## Data Preperation

load the `lalonde.dta`, show the raw mean difference in earning between the control group and the treated group.

```
In [44]: data = lalonde[["treated", "age", "education", "black", "married", "nodegree", "hispanic", "re74", "re75", "re78"]]
data.head()
```

```
Out[44]:
```

	treated	age	education	black	married	nodegree	hispanic	re74	re75	
0	1.0	33	12	0	1	0	0	0.000000	0.000000	124
1	1.0	20	12	0	1	0	0	8644.156250	8644.156250	116
2	0.0	39	12	1	1	0	0	19785.320312	6608.137207	4
3	1.0	49	8	0	1	1	0	9714.596680	7285.947754	167
4	0.0	26	8	0	1	1	0	37211.757812	36941.265625	302

```
In [45]: treatment = data[data.treated == 1]
control = data[data.treated == 0]
```

## Fit Propensity Score Model(s)

First we initialize the `Matcher` object:

- `Matcher` shows the formula used to fit logistic regression model(s) and the number of records in the majority/minority class.
- we use the covariates on the right side of the equation to estimate the probability of being treated (treated = 1).
- Any covariates passed to the (optional) `exclude` parameter will be ignored from the model fitting process. This parameter is particularly useful for unique identifiers like a `userid`.

```
In [46]: m = Matcher(control, treatment, yvar="treated", exclude=["re78"])
```

Formula:

treated ~ age+black+education+hispanic+married+nodegree+re74+re75

n majority: 425

n minority: 297

- The model shows an **imbalance** in our data. The majority group (control group) having more observations than the minority group (treatment group). We account for this by setting `balance=True` in `Matcher.fit_scores()`.
- This tells `Matcher()` to sample from the majority group when fitting the logistic regression model(s) so that the groups are of equal size.
- When undersampling this way, it is highly recommended that `nmodels` is explicitly assigned to a integer much larger than 1. This ensures that more of the majority group is contributing to the generation of propensity scores.
- The value of this integer should depend on the severity of the imbalance: here we use `nmodels=100`.

```
In [47]: np.random.seed(2020)
m.fit_scores(balance=True, nmodels=100)
```

Fitting Models on Balanced Samples: 100\100

Average Accuracy: 55.53%

- The average accuracy of our 100 models is 55.31%, suggesting that there's separability within our data and justifying the need for the matching procedure.
- We don't pay **too much** attention to the estimates of logistic models
- We are interested in **the predicted value of the model** (propensity score of each observation).

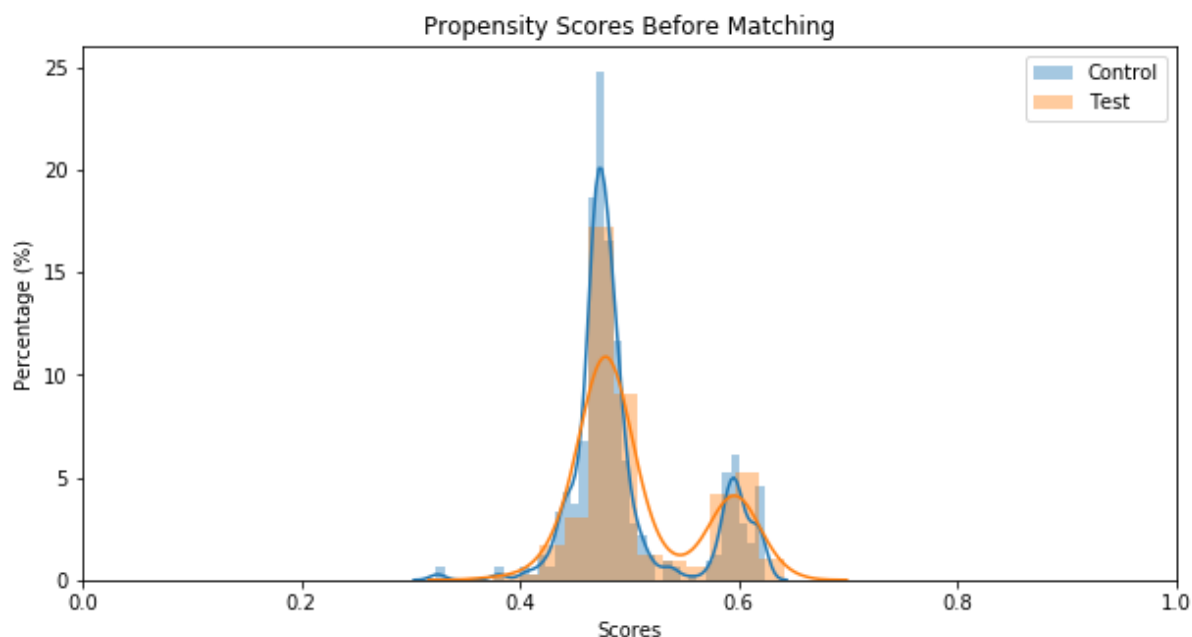
## Assign a propensity score to each record in the dataset

```
In [48]: m.predict_scores()
```

## Evaluating Common Support

We can view the distributions using `m.plot_score()`. It shows that the test and control have good common support.

```
In [49]: m.plot_scores()
```



## K:1 Matching

Then we match one observation from the majority group to each record in the minority group with **replacement**. It means a single majority record can be matched to multiple minority records.

`Matcher` assigns a unique `record_id` to each record in the test and control groups so this can be addressed after matching. If subsequent modeling is planned, one might consider weighting models using a weight vector of  $1/f$  for each record,  $f$  being an observation's frequency in the matched dataset. Thankfully `Matcher` can handle all of this for you :).

```
In [50]: m.match(method="min", nmatches=1) # you can also tune the threshold
```

The matching result is shown as follows: the bulk of our matched-majority-group observation occur only once, 38 occur twice, ... etc.

```
In [51]: m.record_frequency()
```

Out[51]:

	freq	n_records
0	1	422
1	2	38
2	3	21
3	4	7
4	5	1



We can preemptively generate a weight vector using `Matcher.assign_weight_vector()`

```
In [52]: m.assign_weight_vector()
```

## Matched Data

Let's take a look at our matched data!

Note that in addition to the weight vector, `Matcher` has also assigned a `match_id` to each record indicating our (in this case) paired matches since we use `nmatches=1`. We can verify that matched records have scores within 0.0001 of each other.

```
In [53]: matched = m.matched_data.sort_values("match_id")
         matched.head(10)
```

Out[53]:

	record_id	weight	age	black	education	hispanic	married	nodegree	re74	
<b>8</b>	12	0.250000	18	0	12	0	0	0	735.205078	7
<b>297</b>	425	1.000000	33	0	12	0	1	0	0.000000	
<b>9</b>	12	0.250000	18	0	12	0	0	0	735.205078	7
<b>298</b>	426	1.000000	20	0	12	0	1	0	8644.156250	86
<b>106</b>	152	0.250000	19	1	9	0	1	1	0.000000	
<b>299</b>	427	1.000000	49	0	8	0	1	1	9714.596680	72
<b>170</b>	216	0.333333	19	1	12	0	0	0	8417.000000	28
<b>300</b>	428	1.000000	33	1	12	0	1	0	20279.949219	109
<b>93</b>	131	1.000000	25	1	10	0	1	1	0.000000	
<b>301</b>	429	1.000000	35	1	9	0	1	1	13602.429688	138

## Assess Matched Data

```
In [54]: matched.groupby("treated").mean()
```

Out[54]:

	record_id	weight	age	black	education	hispanic	married	nodegree	
<b>treated</b>									
<b>0</b>	203.814815	0.646465	24.292929	0.774411	10.313131	0.111111	0.148148	0.73064	35
<b>1</b>	573.000000	1.000000	24.626263	0.801347	10.380471	0.094276	0.168350	0.73064	35

```
In [55]: # standardized mean difference
def stmean(x):
    x_t = matched[x][matched.treated == 1].mean()
    x_c = matched[x][matched.treated == 0].mean()
    sd_t = matched[x][matched.treated == 1].std()
    stdmean = (x_t-x_c)/sd_t
    print(stdmean)

covariates = ["age", "education", "black", "married", "nodegree", "hispani
c", "re74", "re75"]
for x in covariates:
    stmean(x)

0.049852505643968814
0.03704661050257367
0.06739739852585085
0.053899582655417604
0.0
-0.05751517950116839
0.010446133396393737
0.005471770201405774
```

## Regression with the Matched Data

According to Ho, Imai, King, and Stuart (2007), we can use regression to fit the model.

```
In [56]: # load regression package
import statsmodels.formula.api as smf
import statsmodels.api as sm

weight = matched["weight"].values
smf.wls('re78 ~treated',matched, weights = weight ).fit().summary()
```

Out[56]: WLS Regression Results

<b>Dep. Variable:</b>	re78	<b>R-squared:</b>	0.002
<b>Model:</b>	WLS	<b>Adj. R-squared:</b>	0.001
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1.419
<b>Date:</b>	Mon, 24 Feb 2020	<b>Prob (F-statistic):</b>	0.234
<b>Time:</b>	11:12:02	<b>Log-Likelihood:</b>	-6085.8
<b>No. Observations:</b>	594	<b>AIC:</b>	1.218e+04
<b>Df Residuals:</b>	592	<b>BIC:</b>	1.218e+04
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	5323.2134	427.232	12.460	0.000	4484.138	6162.288
<b>treated</b>	653.1386	548.201	1.191	0.234	-423.517	1729.795

<b>Omnibus:</b>	394.296	<b>Durbin-Watson:</b>	1.933
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	6572.060
<b>Skew:</b>	2.678	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	18.390	<b>Cond. No.</b>	2.95

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [57]: # regression with unmatched data
smf.ols('re78 ~treated', lalonde).fit().summary()
```

Out[57]: OLS Regression Results

<b>Dep. Variable:</b>	re78	<b>R-squared:</b>	0.005
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.003
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	3.525
<b>Date:</b>	Mon, 24 Feb 2020	<b>Prob (F-statistic):</b>	0.0609
<b>Time:</b>	11:12:03	<b>Log-Likelihood:</b>	-7333.1
<b>No. Observations:</b>	722	<b>AIC:</b>	1.467e+04
<b>Df Residuals:</b>	720	<b>BIC:</b>	1.468e+04
<b>Df Model:</b>	1		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	5090.0482	302.783	16.811	0.000	4495.606	5684.491
<b>treated</b>	886.3038	472.086	1.877	0.061	-40.526	1813.134

<b>Omnibus:</b>	384.449	<b>Durbin-Watson:</b>	2.072
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	3767.288
<b>Skew:</b>	2.195	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	13.294	<b>Cond. No.</b>	2.46

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [58]: # you could also add covariates into the regression
smf.wls('re78 ~treated + age+ education + black+ married + nodegree + hispanic + re74 + re75', matched, weights =weight ).fit().summary()
```

Out[58]: WLS Regression Results

<b>Dep. Variable:</b>	re78	<b>R-squared:</b>	0.029
<b>Model:</b>	WLS	<b>Adj. R-squared:</b>	0.014
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	1.925
<b>Date:</b>	Mon, 24 Feb 2020	<b>Prob (F-statistic):</b>	0.0460
<b>Time:</b>	11:12:04	<b>Log-Likelihood:</b>	-6077.8
<b>No. Observations:</b>	594	<b>AIC:</b>	1.218e+04
<b>Df Residuals:</b>	584	<b>BIC:</b>	1.222e+04
<b>Df Model:</b>	9		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	367.5721	2940.969	0.125	0.901	-5408.592	6143.736
<b>treated</b>	606.3762	546.507	1.110	0.268	-466.983	1679.735
<b>age</b>	29.9408	40.656	0.736	0.462	-49.909	109.790
<b>education</b>	398.9888	200.406	1.991	0.047	5.384	792.594
<b>black</b>	-261.3799	864.816	-0.302	0.763	-1959.908	1437.148
<b>married</b>	774.7219	765.501	1.012	0.312	-728.749	2278.193
<b>nodegree</b>	-189.5796	834.935	-0.227	0.820	-1829.420	1450.261
<b>hispanic</b>	858.3063	1164.786	0.737	0.461	-1429.373	3145.986
<b>re74</b>	-0.0135	0.096	-0.140	0.889	-0.202	0.176
<b>re75</b>	0.1005	0.113	0.892	0.373	-0.121	0.322

<b>Omnibus:</b>	397.540	<b>Durbin-Watson:</b>	1.931
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	7054.114
<b>Skew:</b>	2.685	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	19.006	<b>Cond. No.</b>	1.03e+05

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.03e+05. This might indicate that there are strong multicollinearity or other numerical problems.

In [ ]: